

AIX2024 SDK Installation Guide

: Compile & Run **(For *LINUX / MacOS*)**

V1.0

서울대학교 차세대반도체 혁신공유대학

Xuan-Truong Nguyen (응웬트렁)

truongnx@snu.ac.kr

1. Overview

This AIX2024 SDK consists of C/C++ deep learning applications and frameworks based on the darknet framework [1]. This document demonstrates how to install the AIX2024 SDK and associated packages and run applications.

1.1. Directory structure (Next page)

The AIX2024 SDK package is “skeleton” that includes the third-party libraries, source codes, executable folders with the test datasets, weights, and executable scripts, Makefile, and the Visual Studio project.

skeleton

3rdparty	:	Third party library (Don't touch)
lib/		Library
include/		Header files
dll/		Dynamic Link Library
CLBlast/		A modern, lightweight, performant and tunable OpenCL BLAS
library		implements BLAS routines: <i>basic linear algebra subprograms</i> (BLAS) operating on vectors and matrices.
src	:	Source code
main.c		//Main file
yolov2_forward_network.c		// Do inference for an FP32 model
yolov2_forward_network_quantized.c		// Do inference for an int8 quantized
model		
additionally.c		// All functions and utilities
...		
bin	:	Executable files and bash scripts
dataset/		test images and labels, make_list_cur.py, show_images.py, size_search.py, target.txt
weights/		output files when saving the model's parameters layer by layer
*.weights, *.cfg		Files to save the parameters, and the structure of the model
AIX2024		
...		Scripts for Unix/Linux (*.sh), scripts for Windows (*.cmd),
obj	executable files	Object files when compiled on Unix/Linux (Don't care)
Makefile	:	Makefile
*.sln, *.cv*	:	Visual studio project

2. Toolchains for AIX2024 SDK Installation

This chapter walks you through how to set up a development environment, install the AIX2024 SDK, and execute applications on the AIX2024 framework. The following prerequisites and requirements must be satisfied before installing the AIX2024 SDK. We will call these AIX2024 Toolchains. This chapter aims to teach you how to install pre-installed packages for the AIX2024.

The AIX2024 SDK has been tested on the following version of Ubuntu and Python.

- **Ubuntu 18.04.6 LTS (GNU/Linux 4.15.0-142-generic x86_64)**
- **Python version: 3.7.9**
- **GCC version: 7.5.0**

[install required package]

1. Check Python version

```
$ python3 --version  
Python 3.7.9
```

2. Check GCC version

```
$ gcc --version  
gcc (Ubuntu 7.5.0-3ubuntu1~18.04) 7.5.0  
Copyright (C) 2017 Free Software Foundation, Inc.  
This is free software; see the source for copying conditions.  
There is NO warranty; not even for MERCHANTABILITY or  
FITNESS FOR A PARTICULAR PURPOSE.
```

****Note: By the way, the AIX2024 is supposed to work on other versions. If you have any issue related to the package versions, please contact us.**

3. AIX2024 SDK Installation Guide

This chapter elaborates on how to compile the AIX2024 SDK and run a model. Before going to this chapter, you must check Chapter 2 for the required packages. The flow consists of the following steps

- 1) *Generate the directories for the test images (Python)*
- 2) *Compile the code*
- 3) *Execute the scripts*
- 4) *Verify the expected outputs*

Assume that the AIX2024 code is unzipped and stored in your Linux/Unix PC. For example, the AIX2024 is located at `/home/truongnx/aix2024/skeleton/`

1) **Generate the directories for the test images.**

Note that this step is done only **ONE time** when you save the AIX2024 framework in your local directory. Execute the following command lines:

```
$ cd bin/dataset  
$ python3 make_list_cur.py
```

After running those commands, you are supposed to see:

```

CAPP_testset_long_10063
CAPP_testset_close_10031
CAPP_testset_close_10052
CAPP_testset_long_10109
CAPP_testset_long_10007
CAPP_testset_close_10016
CAPP_testset_long_10083
CAPP_testset_close_10075
CAPP_testset_close_10101
CAPP_testset_long_10076
CAPP_testset_long_10042
CAPP_testset_long_10068
CAPP_testset_long_10107
CAPP_testset_close_10073
CAPP_testset_close_10088
CAPP_testset_long_10113
CAPP_testset_close_10062
CAPP_testset_long_10036
(base) truongnx@marlin:~/aix2023/skeleton/bin/dataset$

```

The command executes **make_list_cur.py** to generate all directories of the test images with your local directory and save them to **target.txt**. Note that only jpeg images are included in the file. Now, you can view “target.txt” which stores the directories of all test images.

```
$ vi target.txt
```

```

/home/truongnx/aix2023/skeleton/bin/dataset/CAPP_testset_long_10098.jpg
/home/truongnx/aix2023/skeleton/bin/dataset/CAPP_testset_close_10054.jpg
/home/truongnx/aix2023/skeleton/bin/dataset/CAPP_testset_long_10041.jpg
/home/truongnx/aix2023/skeleton/bin/dataset/CAPP_testset_close_10096.jpg
/home/truongnx/aix2023/skeleton/bin/dataset/CAPP_testset_close_10107.jpg
/home/truongnx/aix2023/skeleton/bin/dataset/CAPP_testset_long_10075.jpg
/home/truongnx/aix2023/skeleton/bin/dataset/CAPP_testset_close_10027.jpg
/home/truongnx/aix2023/skeleton/bin/dataset/CAPP_testset_close_10098.jpg
/home/truongnx/aix2023/skeleton/bin/dataset/CAPP_testset_long_10057.jpg
/home/truongnx/aix2023/skeleton/bin/dataset/CAPP_testset_close_10012.jpg
/home/truongnx/aix2023/skeleton/bin/dataset/CAPP_testset_close_10035.jpg
/home/truongnx/aix2023/skeleton/bin/dataset/CAPP_testset_close_10048.jpg
/home/truongnx/aix2023/skeleton/bin/dataset/CAPP_testset_close_10003.jpg
"target.txt" 229L, 16602C

```

To close the file, you use a combination of keys: “Esc” → “.” → “q” → “!”, which closes the file without any modification.

Now, you can go back to the main folder ***skeleton/*** by using the following command:

```
$ cd ../../
```

2) Compile the AIX2024

****NOTE: Make sure that you are at *skeleton/* before compiling the code.**

- Clean all object files in obj/ and the executable file bin/darknet using the command line:

```
$ make clean
```

- Execute the command line to compile the AIX2024 code:

```
$ make
```

After compilation, you are supposed to see the following screen. And the executable file “darknet” is generated and stored in bin/.

```
gcc -Wall -Wfatal-errors -Ofast -c ./src/yolov2_forward_network_quantized.c -o obj/yolov2_forward_ne
twork_quantized.o
./src/yolov2_forward_network_quantized.c: In function 'gemm_nn_int8_int32':
./src/yolov2_forward_network_quantized.c:11:32: warning: integer overflow in expression [-Woverflow]
#define MAX_VAL_32 (256*256*256*256/2 - 1) // 31-bit (1-bit sign)
                             ^
./src/yolov2_forward_network_quantized.c:104:47: note: in expansion of macro 'MAX_VAL_32'
    C[i*ldc + j] += max_abs(c_tmp[j], MAX_VAL_32);
                                   ^
./src/yolov2_forward_network_quantized.c: In function 'forward_convolutional_layer_q':
./src/yolov2_forward_network_quantized.c:140:0: warning: ignoring #pragma omp parallel [-Wunknown-pra
gmas]
#pragma omp parallel for
./src/yolov2_forward_network_quantized.c: In function 'do_quantization':
./src/yolov2_forward_network_quantized.c:288:26: warning: unused variable 'weights_size' [-Wunused-var
iable]
    size_t const weights_size = l->size*l->size*l->c*l->n;
                             ^
./src/yolov2_forward_network_quantized.c: In function 'save_quantized_model':
./src/yolov2_forward_network_quantized.c:333:26: warning: unused variable 'filter_size' [-Wunused-var
iable]
    size_t const filter_size = l->size*l->size*l->c;
                             ^
In file included from ./src/yolov2_forward_network_quantized.c:1:0:
At top level:
./src/activation.h:130:17: warning: 'activate_array' defined but not used [-Wunused-function]
static void activate_array(float *x, const int n, const ACTIVATION a)
                ^
gcc -Wall -Wfatal-errors -Ofast obj/main.o obj/activation.o obj/box.o obj/yolov2_forward_network.o
obj/yolov2_forward_network_quantized.o -o bin/darknet -lm -pthread
(base) truongnx@marlin:~/aix2023/skeleton$
```

- Now, you can go to bin/ to check if “darknet” is generated.

```
$ cd bin/
$ ll
```

```
(base) truongnx@marlin:~/aix2023/skeleton/bin$ ll
total 10332
drwxrwxr-x 4 truongnx truongnx 4096 15 13:49 ./
drwxrwxr-x 8 truongnx truongnx 4096 15 13:40 ../
-rwxrwxr-x 1 truongnx truongnx 312448 15 13:49 darknet*
drwxrwxr-x 2 truongnx truongnx 28672 15 13:33 dataset/
-rw-rw-r-- 1 truongnx truongnx 2486319 10 13:41 predictions.png
-rw-rw-r-- 1 truongnx truongnx 82944 10 13:41 pthreadVC2.dll
-rw-rw-r-- 1 truongnx truongnx 0 10 13:41 target.txt
-rwxrwxr-x 1 truongnx truongnx 352 10 13:41 tiny-yolo-aix2023-int8.sh*
-rwxrwxr-x 1 truongnx truongnx 112 10 13:41 tiny-yolo-aix2023-int8-test.sh*
-rwxrwxr-x 1 truongnx truongnx 203 10 13:41 tiny-yolo-aix2023.sh*
drwxrwxr-x 2 truongnx truongnx 4096 10 13:41 weights/
-rw-rw-r-- 1 truongnx truongnx 222 10 13:41 yolo_cpu.cmd
-rw-rw-r-- 1 truongnx truongnx 389632 10 13:41 yolo_cpu.exe
-rw-rw-r-- 1 truongnx truongnx 991608 10 13:41 yolo_cpu.ilc
-rw-rw-r-- 1 truongnx truongnx 244 10 13:41 yolo_cpu_int8.cmd
-rw-rw-r-- 1 truongnx truongnx 246 10 13:41 yolo_cpu_int8_test.cmd
-rw-rw-r-- 1 truongnx truongnx 1143878 10 13:41 yolo_cpu.ioobj
-rw-rw-r-- 1 truongnx truongnx 275856 10 13:41 yolo_cpu.ipdb
-rw-rw-r-- 1 truongnx truongnx 733184 10 13:41 yolo_cpu.pdb
-rw-rw-r-- 1 truongnx truongnx 115 10 13:41 yolohw.data
-rw-rw-r-- 1 truongnx truongnx 1323 10 13:41 yolohw.names
-rw-rw-r-- 1 truongnx truongnx 3257 10 13:41 yolov4-tiny-aix2023.cfg
-rw-rw-r-- 1 truongnx truongnx 4059884 10 13:41 yolov4-tiny-aix2023.weights
(base) truongnx@marlin:~/aix2023/skeleton/bin$
```

3) Execute the scripts

****NOTE:** Make sure that you are at **skeleton/bin/** before executing the script.

- Run the **full-precision** model and calculate the mAP by the command:

```
$ script-unix-aix2024-test-all.sh
```

→ This command uses the default “dataset/target.txt” generated at Step 1 and the name list of 60 product items stored in the file “yolohw.names”. Next, it loads the model architecture from aix2024.cfg and then loads the parameters from aix2024.weights.

layer	filters	size	input	output
0 conv	16	3 x 3 / 1	256 x 256 x 3	256 x 256 x 16 0.057 BF
1 max		2 x 2 / 2	256 x 256 x 16	128 x 128 x 16
2 conv	32	3 x 3 / 1	128 x 128 x 16	128 x 128 x 32 0.151 BF
3 max		2 x 2 / 2	128 x 128 x 32	64 x 64 x 32
4 conv	64	3 x 3 / 1	64 x 64 x 32	64 x 64 x 64 0.151 BF
5 max		2 x 2 / 2	64 x 64 x 64	32 x 32 x 64
6 conv	128	3 x 3 / 1	32 x 32 x 64	32 x 32 x 128 0.151 BF
7 max		2 x 2 / 2	32 x 32 x 128	16 x 16 x 128
8 conv	256	3 x 3 / 1	16 x 16 x 128	16 x 16 x 256 0.151 BF
9 max		2 x 2 / 2	16 x 16 x 256	8 x 8 x 256
10 conv	512	3 x 3 / 1	8 x 8 x 256	8 x 8 x 512 0.151 BF
11 max		2 x 2 / 1	8 x 8 x 512	8 x 8 x 512
12 conv	256	1 x 1 / 1	8 x 8 x 512	8 x 8 x 256 0.017 BF
13 conv	512	3 x 3 / 1	8 x 8 x 256	8 x 8 x 512 0.151 BF
14 conv	195	1 x 1 / 1	8 x 8 x 512	8 x 8 x 195 0.013 BF
15 yolo				
16 route	12			
17 conv	128	1 x 1 / 1	8 x 8 x 256	8 x 8 x 128 0.004 BF
18 upsample		2x	8 x 8 x 128	16 x 16 x 128
19 route	18 8			
20 conv	195	1 x 1 / 1	16 x 16 x 384	16 x 16 x 195 0.038 BF
21 yolo				
Total	BFLOPS	1.035		

4) Verify the expected outputs

Finally, it executes the model on 229 test images and then calculates the mAP. You are supposed to see **“mean average precision (mAP) = 0.817559, or 81.76 %”**.

Depending on your PC specifications, it may take more or less execution.

```

class_id = 24, name = reeses_pieces, ap = 100.00 %
class_id = 25, name = clif_crunch_peanut_butter, ap = 73.53 %
class_id = 26, name = mom_to_mom_butternut_squash_pear,
ap = 90.05 %
class_id = 27, name = pop_tararts_strawberry, ap = 91.21 %
class_id = 28, name = quaker_big_chewy_chocolate_chip, ap = 77.93 %
class_id = 29, name = spam, ap = 61.59 %
class_id = 30, name = coffee_mate_french_vanilla, ap = 57.76 %
class_id = 31, name = pepperidge_farm_milk_chocolate_macadamia_cookies, ap = 74.59 %
class_id = 32, name = kitkat_king_size, ap = 60.20 %
class_id = 33, name = snickers, ap = 11.17 %
class_id = 34, name = toblerone_milk_chocolate, ap = 41.16 %
class_id = 35, name = clif_z_bar_chocolate_chip, ap = 97.71 %
class_id = 36, name = nature_valley_crunchy_oats_n_honey, ap = 86.78 %
class_id = 37, name = ritz_crackers, ap = 100.00 %
class_id = 38, name = palmolive_orange, ap = 87.18 %
class_id = 39, name = crystal_hot_sauce, ap = 85.22 %
class_id = 40, name = tapatio_hot_sauce, ap = 66.87 %
class_id = 41, name = nabisco_nilla_wafers, ap = 85.76 %
class_id = 42, name = pepperidge_farm_milano_cookies_double_chocolate, ap = 94.20 %
class_id = 43, name = campbells_chicken_noodle_soup, ap = 99.47 %
class_id = 44, name = frappuccino_coffee, ap = 91.53 %
class_id = 45, name = chewy_dips_chocolate_chip, ap = 64.94 %
class_id = 46, name = chewy_dips_peanut_butter, ap = 89.97 %
class_id = 47, name = nature_vally_fruit_and_nut, ap = 92.30 %
class_id = 48, name = cheerios, ap = 96.27 %
class_id = 49, name = lindt_excellence_cocoa_dark_chocolate, ap = 81.82 %
class_id = 50, name = hersheys_symphony, ap = 100.00 %
class_id = 51, name = campbells_chunky_classic_chicken_noodle, ap = 94.65 %
class_id = 52, name = martinellis_apple_juice, ap = 79.72 %
class_id = 53, name = dove_pink, ap = 74.48 %
class_id = 54, name = dove_white, ap = 88.93 %
class_id = 55, name = david_sunflower_seeds, ap = 95.94 %
class_id = 56, name = monster_energy, ap = 44.72 %
class_id = 57, name = act_ii_butter_lovers_popcorn, ap = 86.10 %
class_id = 58, name = coca_cola_glass_bottle, ap = 81.61 %
class_id = 59, name = twix, ap = 85.90 %
for thresh = 0.24, precision = 0.80, recall = 0.70, F1-score = 0.74
for thresh = 0.24, TP = 2058, FP = 508, FN = 901, average IoU = 59.16 %

mean average precision (mAP) = 0.817559, or 81.76 %
Total Detection Time: 36.000000 Seconds
bo@BoGram14:/mnt/c/skeleton/bin$

```

- Run the **int8 quantized** model and calculate the mAP by the command:

```
$ script-unix-aix2024-test-all-quantized.sh
```

→ Now, you can see some similar outputs as that of the full-precision model.

However, after loading the model, it prints out the default quantization multipliers for an input image or input feature maps, weights, and biases.

layer	filters	size	input	output
0 conv	16	3 x 3 / 1	256 x 256 x 3	256 x 256 x 16 0.057 BF
1 max		2 x 2 / 2	256 x 256 x 16	128 x 128 x 16
2 conv	32	3 x 3 / 1	128 x 128 x 16	128 x 128 x 32 0.151 BF
3 max		2 x 2 / 2	128 x 128 x 32	64 x 64 x 32
4 conv	64	3 x 3 / 1	64 x 64 x 32	64 x 64 x 64 0.151 BF
5 max		2 x 2 / 2	64 x 64 x 64	32 x 32 x 64
6 conv	128	3 x 3 / 1	32 x 32 x 64	32 x 32 x 128 0.151 BF
7 max		2 x 2 / 2	32 x 32 x 128	16 x 16 x 128
8 conv	256	3 x 3 / 1	16 x 16 x 128	16 x 16 x 256 0.151 BF
9 max		2 x 2 / 2	16 x 16 x 256	8 x 8 x 256
10 conv	512	3 x 3 / 1	8 x 8 x 256	8 x 8 x 512 0.151 BF
11 max		2 x 2 / 1	8 x 8 x 512	8 x 8 x 512
12 conv	256	1 x 1 / 1	8 x 8 x 512	8 x 8 x 256 0.017 BF
13 conv	512	3 x 3 / 1	8 x 8 x 256	8 x 8 x 512 0.151 BF
14 conv	195	1 x 1 / 1	8 x 8 x 512	8 x 8 x 195 0.013 BF
15 yolo				
16 route	12			
17 conv	128	1 x 1 / 1	8 x 8 x 256	8 x 8 x 128 0.004 BF
18 upsample		2x	8 x 8 x 128	16 x 16 x 128
19 route	18 8			
20 conv	195	1 x 1 / 1	16 x 16 x 384	16 x 16 x 195 0.038 BF
21 yolo				

Total BFLOPS 1.035
Loading weights from aix2024.weights...
Done!

Multiplier	Input	Weight	Bias
CONV0:	128	16	2048
CONV2:	16	64	1024
CONV4:	16	64	1024
CONV6:	16	64	1024
CONV8:	16	64	1024
CONV10:	16	64	1024
CONV12:	16	64	1024
CONV13:	16	64	1024
CONV14:	16	64	1024
CONV17:	16	64	1024
CONV20:	16	64	1024

Finally, it executes the int8 quantized model on 229 test images and then calculates the mAP. You are supposed to see “**mean average precision (mAP) = 0.550382, or 55.04 %**”. Depending on your PC specifications, it may take more or less execution.

Since we used the default multiplier for quantization, the mAP result is currently poor. It's your job to improve it by following the 'Quantization Manual.pdf' and Tutorial 03 on Quantization."

```
class_id = 38, name = palmolive_orange, ap = 69.79 %
class_id = 39, name = crystal_hot_sauce, ap = 32.63 %
class_id = 40, name = tapatio_hot_sauce, ap = 61.68 %
class_id = 41, name = nabisco_nilla_wafers, ap = 89.12 %
class_id = 42, name = pepperidge_farm_milano_cookies_double_chocolate, ap = 73.35 %
class_id = 43, name = campbells_chicken_noodle_soup, ap = 45.40 %
class_id = 44, name = frappuccino_coffee, ap = 80.00 %
class_id = 45, name = chewy_dips_chocolate_chip, ap = 35.23 %
class_id = 46, name = chewy_dips_peanut_butter, ap = 72.98 %
class_id = 47, name = nature_vally_fruit_and_nut, ap = 24.98 %
class_id = 48, name = cheerios, ap = 87.83 %
class_id = 49, name = lindt_excellence_cocoa_dark_chocolate, ap = 59.94 %
class_id = 50, name = hersheys_symphony, ap = 98.99 %
class_id = 51, name = campbells_chunky_classic_chicken_noodle, ap = 66.98 %
class_id = 52, name = martinellis_apple_juice, ap = 32.10 %
class_id = 53, name = dove_pink, ap = 22.91 %
class_id = 54, name = dove_white, ap = 58.06 %
class_id = 55, name = david_sunflower_seeds, ap = 71.00 %
class_id = 56, name = monster_energy, ap = 10.08 %
class_id = 57, name = act_ii_butter_lovers_popcorn, ap = 47.84 %
class_id = 58, name = coca_cola_glass_bottle, ap = 65.94 %
class_id = 59, name = twix, ap = 37.09 %
for thresh = 0.24, precision = 0.64, recall = 0.23, F1-score = 0.34
for thresh = 0.24, TP = 672, FP = 380, FN = 2287, average IoU = 44.83 %

mean average precision (mAP) = 0.550382, or 55.04 %
Total Detection Time: 9.000000 Seconds
```

You can also run a script to test the AIX 2024 model on one image

```
$ script-unix-aix2024-test-one.sh
```

```
$ script-unix-aix2024-test-one-quantized.sh
```

References

- [1]. <https://github.com/pjreddie/darknet>