

Quantization

Manual

AIX 2024

서울대학교 차세대반도체 혁신공유대학

Xuan-Truong Nguyen (응웬트렁)

truongnx@snu.ac.kr

If you have no background in **YOLO** or **Quantization**, please refer to the links below:

YOLO:

<https://www.notion.so/aix2024tutorial02/02-YOLO-network-AIX-tutorial-3b4314c624bf4007a4268df33b0d6b82>

Quantization:

<https://www.notion.so/aix2024tutorial02/03-Quantization-AIX-tutorial-5a03d7366197428cbf1454e4be97057c>

1. Source code related to quantization

(in C:\skeleton\src)

- **yolov2_forward_network_quantized.c**
= Functions for quantization, saving of the quantized model, and the forward pass of quantized yolo model
(↑ You should mainly edit this file for quantization!)
- main.c // *The main functions*

In “main.c”, there is a function named “test_detector_cpu”.

[test_detector_cpu] □ main.c

```
void test_detector_cpu(  
    char **names,           // List of all items (bin/yolohw.names)  
    char *cfgfile,          // Configuration file (bin/aix2024.cfg)  
    char *weightfile,       // Configuration file (bin/aix2024.weights)
```

```

char *filename,           // Input image file
float thresh,             // Hierarchical threshold
int quantized,            // On/off quantization
int save_params,          // On/off save output
int dont_show             // Don't show
)

```

This function **parse the configuration file** and **call an inference function**.

```

if (quantized) {
    network_predict_quantized(net, x);    // quantized
    nms = 0.2;
}
else {
    network_predict_cpu(net, x);
}

```

A network **architecture** is stored in the variable “**net**”.

```

network net = parse_network_cfg(cfgfile, 1, quantized);    // parser.c

```

2. Weight/bias/activation Quantization

The code below represents a series of operations that the accelerator should perform for detection. Whether you use a full precision model or a quantized model, detection goes through the following process.

```

137 // Use GEMM (as part of BLAS)
138 im2col_cpu_int8(state.input_int8, l.c, l.h, l.w, l.size, l.stride, l.pad, b);
139 int m; // multi-thread gemm
140 #pragma omp parallel for
141 for (t = 0; t < m; ++t) {
142     gemm_nn_int8_int16(1, n, k, 1, a + t*k, k, b, n, c + t*n, n);
143 }
144 free(state.input_int8);
145
146 // Bias addition
147 int fil;
148 for (fil = 0; fil < l.n; ++fil) {
149     for (j = 0; j < out_size; ++j) {
150         output_q[fil*out_size + j] = output_q[fil*out_size + j] + l.biases_quant[fil];
151     }
152 }
153
154 // Activation
155 if (l.activation == LEAKY) {
156     for (i = 0; i < l.n*out_size; ++i) {
157         output_q[i] = (output_q[i] > 0) ? output_q[i] : output_q[i] / 10;
158     }
159 }
160
161 // De-scaling
162 float ALPHA1 = 1 / (l.inp);
163 for (i = 0; i < l.outputs; ++i) {
164     l.output[i] = output_q[i] * ALPHA1;
165 }
166 free(output_q);
167 }
168 }

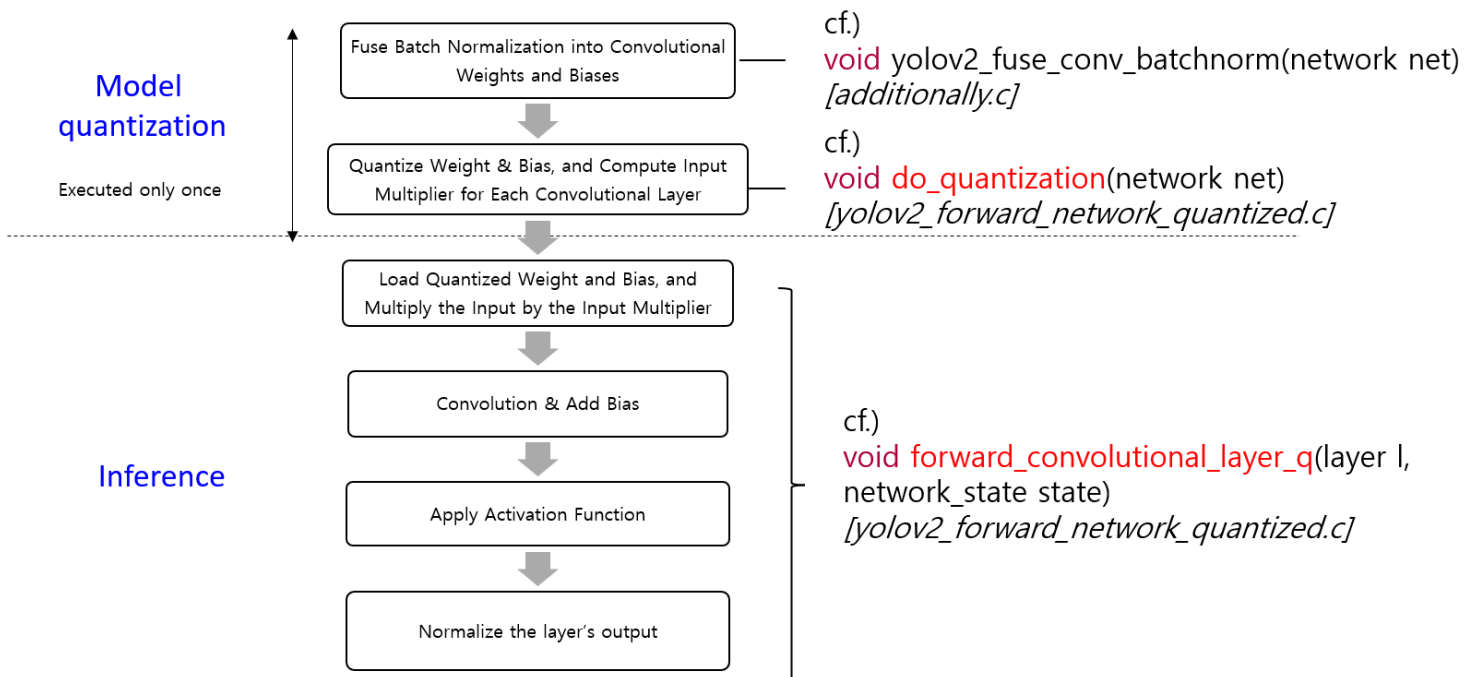
```

The accelerator should perform these operations

- im2col prepares the input image for convolution by converting it into a matrix.
- GEMM performs the convolution operation by multiplying the input matrix by the kernel matrix.

These two operations are essential for efficient implementation of convolutions in deep learning and are often used together.

The difference between the two is determined by whether or not the data is quantized prior to detection. The process of using a quantized model is as follows.



Let's see how quantization is implemented in the code.

[do_quantization] □ @ yolov2_forward_network_quantized.c

- Input:
 - Network
- Output
 - Quantized model

In AIX 2024, it is our job to **find best set of multipliers** □

for quantization. It is now set to default values.

```
129 void do_quantization(network net) {
130     int counter = 0;
131
132     int j;
133     //Dummy weight quantization
134     #define TOTAL_CALIB_LAYER 18
135     //
136     float input_quant_multiplier[TOTAL_CALIB_LAYER] = {
137         128, //conv 0
138         8, //conv 1
139         8, //conv 2
140         8, //conv 4
141         8, //conv 5
142         8, //conv 7
143         8, //conv 10
144         8, //conv 12
145         8, //conv 13
146         8, //conv 15
147         8, //conv 18
148         8, //conv 20
149         8, //conv 21
150         8, //conv 23
151         8, //conv 26
152         8, //conv 27
153         8, //conv 30
154         8, //conv 33
155     };
156
157     float weight_quant_multiplier[TOTAL_CALIB_LAYER] = {
158         16, //conv 0
159         128, //conv 1
160         128, //conv 2
161         128, //conv 4
162         128, //conv 5
163         128, //conv 7
164         128, //conv 10
165         128, //conv 12
166         128, //conv 13
167         128, //conv 15
168         128, //conv 18
169         128, //conv 20
170         128, //conv 21
171         128, //conv 23
172         128, //conv 26
173         128, //conv 27
174         128, //conv 30
175         128, //conv 33
176     };
```

As described in the tutorial03, the goal of quantization is to convert **from FP32 to a fixed-point format**. For AIX 2024 quantization, **we use 8 bits for weights and 16 bits for biases**.

- Two steps to quantize **weights and biases**
 - 1) **Scaling**: Multiply by a scale factor(multiplier)
 - 2) **Clipping**: avoid overflow

Ex) int8 □ if $x > 127$, $x = 127$ or if $x < -128$, $x = -128$

For example, **multiply a given weight by 16 and apply the round-to-nearest on the result**; we can obtain a fixed-point number with four fractional bits, three integer bits, and one sign bit.

↓ You can also **implement your own quantization** to minimize mAP degradation.

```
277 printf("Multiplier Input weight bias\n");
278 for (j = 0; j < net.n; ++j) {
279     layer *l = &net.layers[j];
280
281     /* ... */
282
283     //printf("\n");
284     if (l->type == CONVOLUTIONAL) { // Quantize conv layer only
285         size_t const weights_size = l->size*l->size*l->nc*l->no;
286         size_t const filter_size = l->size*l->size*l->nc;
287
288         int i, fil;
289
290         // scaling
291         //{{{
292         // Input feature map
293         l->input_quant_multiplier = (counter < TOTAL_CALIB_LAYER) ? input_quant_multiplier[counter] : 16;
294
295         // weight
296         l->weights_quant_multiplier = (counter < TOTAL_CALIB_LAYER) ? weight_quant_multiplier[counter] : 16;
297
298         ++counter;
299         //}}}
300         for (fil = 0; fil < l->no; ++fil) {
301             for (i = 0; i < filter_size; ++i) {
302                 float w = l->weights[fil*filter_size + i] * l->weights_quant_multiplier; // scale
303                 l->weights_ints[fil*filter_size + i] = max_abs(w, MAX_VAL_8); // clip
304             }
305         }
306
307         // Bias Quantization
308         float biases_multiplier = (l->weights_quant_multiplier * l->input_quant_multiplier);
309         for (fil = 0; fil < l->no; ++fil) {
310             float b = l->biases[fil] * biases_multiplier; // scale
311             l->biases_quant[fil] = max_abs(b, MAX_VAL_16); // clip
312         }
313
314         //printf("conv%d multipliers: input %g, weights %g, bias %g\n", j, l->input_quant_multiplier, l->weights_quant_multiplier, biases_multiplier);
315         printf(" conv%d: %tq %tq %tq\n", j, l->input_quant_multiplier, l->weights_quant_multiplier, biases_multiplier);
316     }
317     else {
318         //printf("No quantization for layer %d (layer type: %d)\n", j, l->type);
319     }
320 }
```

[forward_convolutional_layer_q] □ @ yolov2_forward_network_quantized.c

- Similar to weight, **activation** can be quantized by multiplying it by a factor, i.e., input_quant_multiplier, and applying the round-to-nearest operation on the multiplication result.
- Again, activation quantization also aims to find a fixed-point format of a floating-point number. More importantly, **it allows to do a 8-bit x 8-bit multiplication** (gemm_nn_int8_int16 at Line 142).

3. Descaling

After convolution, it adds a bias, performs activation, and then do **descaling**.

Descaling can be considered as conversion from a fixed-point number to a floating-point one. For the model's **outputs to be interpretable**, they must be **converted back to their original scale**.

```
145
146 // Bias addition
147 int fil;
148 for (fil = 0; fil < 1.n; ++fil) {
149     for (j = 0; j < out_size; ++j) {
150         output_q[fil*out_size + j] = output_q[fil*out_size + j] + 1.biases_quant[fil];
151     }
152 }
153
154 // Activation
155 if (1.activation == LEAKY) {
156     for (i = 0; i < 1.n*out_size; ++i) {
157         output_q[i] = (output_q[i] > 0) ? output_q[i] : output_q[i] / 10;
158     }
159 }
160
161 // De-scaling
162 float ALPHA1 = 1 / (1.input_quant_multiplier * 1.weights_quant_multiplier);
163 for (i = 0; i < 1.outputs; ++i) {
164     1.output[i] = output_q[i] * ALPHA1;
165 }
166
167 free(output_q);
168 }
```

We can **merge a descaling step to an input quantization** (Lines 125-126). In other words, we want to store input_int8.

```

111 void forward_convolutional_layer_q(layer l, network_state state)
112 {
113
114     int out_h = (l.h + 2 * l.pad - l.size) / l.stride + 1; // output_height=input_height for stride=1 and pad=1
115     int out_w = (l.w + 2 * l.pad - l.size) / l.stride + 1; // output_width=input_width for stride=1 and pad=1
116     int i, j;
117     int const out_size = out_h*out_w;
118
119     typedef int16_t conv_t; // l.output
120     conv_t *output_q = calloc(l.outputs, sizeof(conv_t));
121
122     state.input_int8 = (int8_t *)calloc(l.inputs, sizeof(int));
123     int z;
124     for (z = 0; z < l.inputs; ++z) {
125         int16_t src = state.input[z] * l.input_quant_multiplier;
126         state.input_int8[z] = max_abs(src, MAX_VAL_8);
127     }
128
129     // Convolution
130     int m = l.n;
131     int k = l.size*l.size*l.c;
132     int n = out_h*out_w;
133     int8_t *a = l.weights_int8;
134     int8_t *b = (int8_t *)state.workspace;
135     conv_t *c = output_q; // int16_t
136
137     // Use GEMM (as part of BLAS)
138     im2col_cpu_int8(state.input_int8, l.c, l.h, l.w, l.size, l.stride, l.pad, b);
139     int t; // multi-thread gemm
140     #pragma omp parallel for
141     for (t = 0; t < m; ++t) {
142         gemm_nn_int8_int16(1, n, k, 1, a + t*k, k, b, n, c + t*n, n);
143     }
144     free(state.input_int8);

```


[yolov2_forward_network_q] @ yolov2_forward_network_quantized.c

This function is similar to yolov2_forward_network, which **takes a model and an image and generates the detected tensors**. It is also executed layer by layer, where a layer can be convolution, max-pooling, route, or upsample.

```
170 void yolov2_forward_network_q(network net, network_state state)
171 {
172     state.workspace = net.workspace;
173     int i;
174     for (i = 0; i < net.n; ++i) {
175         state.index = i;
176         layer l = net.layers[i];
177
178         if (l.type == CONVOLUTIONAL) {
179             forward_convolutional_layer_q(l, state);
180         }
181         else if (l.type == MAXPOOL) {
182             forward_maxpool_layer_cpu(l, state);
183         }
184         else if (l.type == ROUTE) {
185             forward_route_layer_cpu(l, state);
186         }
187         else if (l.type == REORG) {
188             forward_reorg_layer_cpu(l, state);
189         }
190         else if (l.type == UPSAMPLE) {
191             forward_upsample_layer_cpu(l, state);
192         }
193         else if (l.type == SHORTCUT) {
194             forward_shortcut_layer_cpu(l, state);
195         }
196         else if (l.type == YOLO) {
197             forward_yolo_layer_cpu(l, state);
198         }
199         else if (l.type == REGION) {
200             forward_region_layer_cpu(l, state);
201         }
202         else {
203             printf("\n layer: %d \n", l.type);
204         }
205         state.input = l.output;
206     }
207 }
```

Note that the **final output** consists of the two tensors **8x8x195** and **16x16x195** from **Layers 27 and 33**, respectively, which are **post-processed by the YOLO layers to generate bounding boxes**.

4. Prepare Quantized Data

1) Quantized Model

`[save_quantized_model]` □ `@ yolov2_forward_network_quantized.c`

You can use a template provided to store weights and biases layer by layer. The quantized model is stored in your computing order. Our quantized weights are generally stored in DRAM or block ram (BRAM). From DRAM or BRAM, a CNN engine must load weights into registers, when computing each layer.

```

326 // Save quantized weights, bias, and scale
327 void save_quantized_model(network net) {
328     int j;
329     for (j = 0; j < net.n; ++j) {
330         layer *l = &net.layers[j];
331         if (l->type == CONVOLUTIONAL) {
332             size_t const weights_size = l->size*l->size*l->c*l->n;
333             size_t const filter_size = l->size*l->size*l->c;
334
335             printf(" Saving quantized weights, bias, and scale for CONV%d \n", j);
336
337             char weightfile[30];
338             char biasfile[30];
339             char scalefile[30];
340
341             sprintf(weightfile, "weights/CONV%d_W.txt", j);
342             sprintf(biasfile, "weights/CONV%d_B.txt", j);
343             sprintf(scalefile, "weights/CONV%d_S.txt", j);
344
345             int k;
346
347             FILE *fp_w = fopen(weightfile, "w");
348             for (k = 0; k < weights_size; k = k + 4) {
349                 uint8_t first = k < weights_size ? l->weights_int8[k] : 0;
350                 uint8_t second = k+1 < weights_size ? l->weights_int8[k+1] : 0;
351                 uint8_t third = k+2 < weights_size ? l->weights_int8[k+2] : 0;
352                 uint8_t fourth = k+3 < weights_size ? l->weights_int8[k+3] : 0;
353                 fprintf(fp_w, "%02x%02x%02x%02x\n", first, second, third, fourth);
354             }
355             fclose(fp_w);
356
357             FILE *fp_b = fopen(biasfile, "w");
358             for (k = 0; k < l->n; k = k + 4) {
359                 uint16_t first = k < l->n ? l->biases_quant[k] : 0;
360                 uint16_t second = k+1 < l->n ? l->biases_quant[k+1] : 0;
361                 fprintf(fp_b, "%04x%04x\n", first, second);
362             }
363             fclose(fp_b);

```

It opens a file for the weights and writes the quantized weights(8bit each) to the file in hexadecimal format, four weights per line.

It opens a file for the biases and writes the quantized biases(16bit each) to the file in hexadecimal format, two biases per line.

2) Quantized Input image

- **Input image:** The input image has a size of 256x256x3, where pixels are ordered in a 1D array. Note that the **input image is normalized to [0, 1]**. When executing the first layer, the input image is **multiplied by an input quantized multiplier** to generate fixed-point pixels.

⇒ Consider saving input in int8 format and sending it to DRAM.

- **Input feature maps:** Each convolution layer may take the output of its previous layer as an input. Therefore, it is necessary to **store input_int8 in a specific format as a test vector.**

3) Output feature maps

- To verify the functionality of a CNN engine, it requires **generating a test vector for output.** When designing an engine, we must compare the output from the HW simulation and the output from the AIX2024 SDK. Make sure that the results are identical. Again, you could use the sample template to write a file.

References

[4]. Jacob, “Quantization and Training of Neural Networks for Efficient Integer-Arithmetic-Only Inference,” CVPR 2018.