

DSP Manual

With Nexys A7-100T

서울대학교 차세대반도체 혁신공유대학

Xuan-Truong Nguyen (응웬트렁)

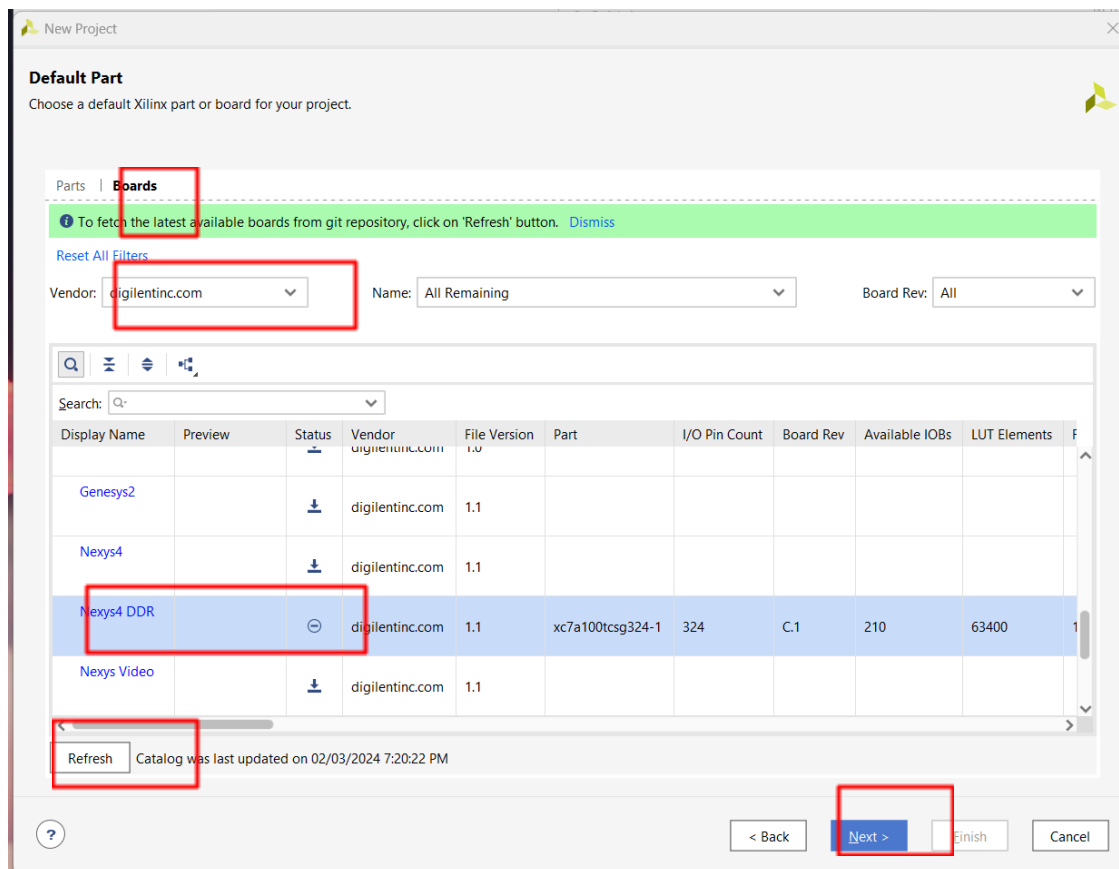
truongnx@snu.ac.kr

1. Nexys A7-100T in Vivado

To examine the board Nexys A7-100T in Vivado, you can do the following steps:

- **Open a new project** with Vivado
- Click **Next** until it opens the window “**Default Part**”, “Choose a default Xilinx part or board for your project”
- Select Tab “**Boards**” → Click “**Refresh**” □ Select **Vendor** “diligentinc.com” → Select “**Nexys4 DDR**” □ Install

⇒ You can examine many parameters of the FPGA chip xc7a100tcsq324-1 like 324 IO pins, 210 IOBs, 63,400 (=15,850×4) LUTs, 126,800 (=15,850×8) FFs, 135 BRAMs (each BRAM = 36kbit), and 240 DSPs.



- Now, you can select Nexys4 DDR(or Nexys A7-100T in the later version Vivado, since the Nexys 4 DDR has been replaced by the Nexys A7)

** When you actually use the board later, you can download&add the external IP and use the same 'Nexys A7-100T' board in the old version Vivado too.*

- Click Next to create a new project.

Project Summary

Overview | Dashboard

Board revision:

D.0

Connectors:

No connections

Repository path:

E:/Xilinx/Vivado/2021.1/data/boards/board_files

URL:

<https://reference.digilentinc.com/reference/programmable-logic/nexys-a7/start>

Board overview:

Nexys A7-100T

Synthesis

Status:

Not started

Messages:

No errors or warnings

Part:

xc7a100tcs9324-1

Strategy:

[Vivado Synthesis Defaults](#)

Report Strategy:

[Vivado Synthesis Default Reports](#)

Incremental synthesis:

None

DRC Violations

[Run Implementation](#) to see DRC results

Utilization

[Run Synthesis](#) to see utilization results

Implementation

Status:

Not started

Messages:

No errors or warnings

Part:

xc7a100tcs9324-1

Strategy:

[Vivado Implementation Defaults](#)

Report Strategy:

[Vivado Implementation Default Reports](#)

Incremental implementation:

None

Timing

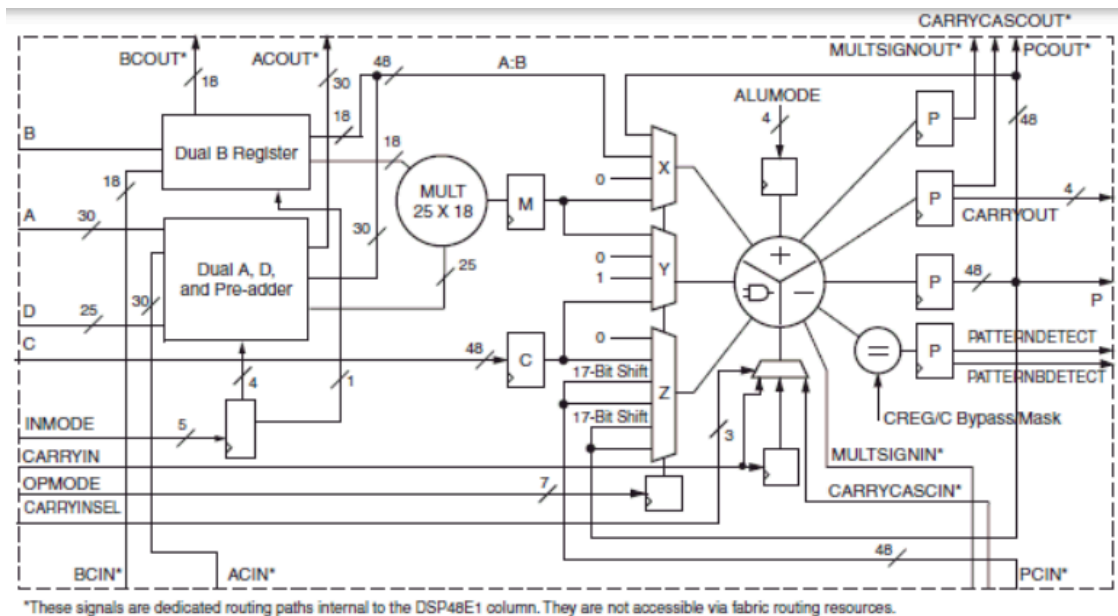
[Run Implementation](#) to see timing results

Power

[Run Implementation](#) to see power results

2. Computing units and DSP

As described in the Tutorial 05, the primary operations in the CNN network are **multiplication and accumulation (MAC)** that can be mapped to the **pre-built DSP** on an **FPGA board [1]**. The DSP block is built on a **25 x 18 bit multiplier** (for details see [1]).

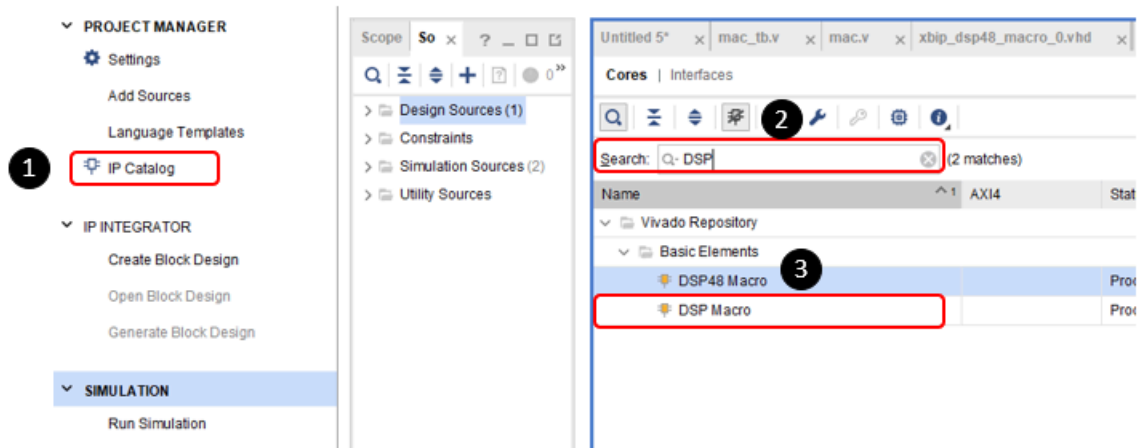


Let's implement a multiplication based on DSP. To do this, you need to...

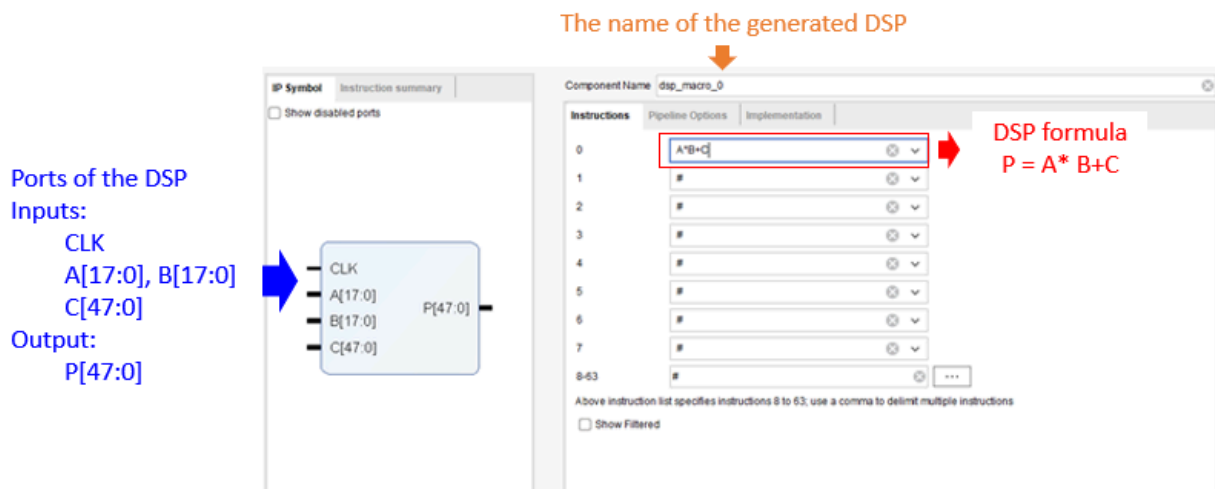
- 1) **Build a DSP using Xilinx IP generator**
- 2) **Use(instantiate) the pre-built DSP to implement a multiplier**

1) Build a DSP using Xilinx IP generator

First, you can use “IP Catalog” and find “DSP Macro”.



Select **Instructions**, for example, $A * B + C$



Configure **Pipeline Options** for a DSP block.

- Note that you can change **Automatic** into **Expert** to remove some pipeline registers.

- For example, the default DSP takes **four cycles to complete a MAC**. However, if you uncheck all pipeline registers, it takes zero-cycle to execute MAC as a combinational circuit.

The screenshot shows the configuration interface for a DSP macro. On the left is the IP Symbol block with ports CLK, A[17:0], B[17:0], C[47:0], and output P[47:0]. The main panel is titled 'Component Name dsp_macro_0' and has tabs for Instructions, Pipeline Options, and Implementation. The Pipeline Options tab is active, showing 'Pipeline Options' set to 'Automatic'. Below this is a 'Custom Pipeline options' table with 6 tiers. The table shows the internal logic flow: input A is added to B in tier 3, the result is multiplied by C in tier 4, and the final result is concatenated with C in tier 5. Blue arrows point to the pipeline registers between tiers 1-2, 2-3, 3-4, and 4-5, with a note: 'Pipeline registers (4-cycle delay between an input and an output)'. A key at the bottom indicates that yellow cells represent 'Fabric register' and white cells represent 'DSP register'.

Tier:	1	2	3	4	5	6
	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>
D	<input type="checkbox"/>	→	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>
A	<input type="checkbox"/>	→	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>
B	<input type="checkbox"/>	→	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>
CONCAT	→		<input type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>
C	<input type="checkbox"/>	→	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>
CARRYIN	<input type="checkbox"/>	→	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>
SEL	<input type="checkbox"/>	→	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>
KEY:	Fabric register					
	DSP register					

Step 4) Configure **Implementation** for a DSP block.

You can change the size of inputs and outputs. Although it **does not change the DSP cost**, it saves some resources of the external module that uses a DSP because it works with fewer bits.

IP Symbol

Instruction summary

☐ Show disabled ports

Component Name

dsp_macro_0

Instructions

Pipeline Options

Implementation

Input Port Properties

D

18

[2 - 25]

A

18

[2 - 25]

B

18

[2 - 18]

C

48

[2 - 48]

CONCAT

48

[2 - 48]

Output Port Properties

Output Properties

Full Precision

Full Precision Width 48

Width

48

[2 - 48]

Additional ports

☐ Use ACOUT
☐ Use BCOUT
☐ Use PCOUT

☐ Use CARRYOUT
☐ Use CARRYCASOUT

☒ Use DSP Slice

Step 5) Click OK and “Generate” on the popup window

<DSP Generated>

Tcl ConsoleMessagesLogReportsDesign Runs

Q

≡

⏏

⏮

⏪

⏩

⏭

+

%

Name	Constraints	Status	WNS	TNS	WHS	THS	TPWS	Total Power	Failed Ports	LUT	FF	BRAM	URAM	DSP	Start	Elapsed	Run Strategy
synth_1 (active)	constrs_1	Not started															Vivado Synthesis Defaults (V
impl_1	constrs_1	Not started															Vivado Implementation Defaul
Out-of-Context Module Runs																	
dsp_macro_0_synth_1	dsp_macro_0	synth_design Complete!								49	146	0.0	0	1	2/1/24	00:00:28	Vivado Synthesis Defaults (Viv

Hardware resource

One DSP is used

2) Use the pre-built DSP to implement a multiplier

Now that we generated our DSP, we can use it by 'instantiating' the module in verilog code.

Before going to the instantiation, note that you can implement a multiplier either (1) with prebuilt DSP or (2) with LUTs.

```
1  `timescale 1ns / 1ps
2
3  module mul_dsp(
4      clk,
5      w_i,
6      x_i,
7      y_o
8  );
9      parameter DATA_WIDTH = 16;
10     input clk;
11     input signed [ DATA_WIDTH -1 :0] w_i;
12     input signed [ DATA_WIDTH -1 :0] x_i;
13     output signed [2 * DATA_WIDTH -1 :0] y_o;
14
15     // Combinational clk
16     assign y_o = x_i * w_i;
17
18 endmodule
19
```

(1) Multiplier with DSP

```
1  `timescale 1ns / 1ps
2  (*use_dsp48 = "no"*)
3  module mul_lut(
4      clk,
5      w_i,
6      x_i,
7      y_o
8  );
9      parameter DATA_WIDTH = 16;
10     input clk;
11     input signed [ DATA_WIDTH -1 :0] w_i;
12     input signed [ DATA_WIDTH -1 :0] x_i;
13     output signed [2 * DATA_WIDTH -1 :0] y_o;
14
15     // Combinational clk
16     assign y_o = x_i * w_i;
17
18 endmodule
19
```

(2) multiplier with LUT

The following codes are two multiplier designs, one using DSP and another using LUT.

To force Vivado to implement a multiplier **without using DSP**, you can use the flag **(*use_dsp48="no"*)** as shown in Line 2 of Version (b).

Without using the flag, if DATA_WIDTH is less than or equal to 10, a multiplier is **automatically mapped to LUTs**. Meanwhile, if DATA_WIDTH is greater or equal to 11, **a multiplier is mapped to DSPs and LUTs**.

Once you choose to use multiplier based on DSP, you can manually **instantiate** the DSP modules to use. When you want to utilize DSP for your eight-bit multiplier, you can **call a DSP instance** as follows:

```
dsp_macro_0 u_dsp(
  .CLK(clk),
  .A(dsp_A),
  .B(dsp_B),
  .C(48'b0),
  .P(dsp_P)
);
```

The following table shows the resources for different DATA_WIDTHs. If you implement 240 multipliers without using DSPs, it will cost 14,640 LUTs accounting for 23.09% of the total LUTs in a Nexys A7-100T board. **In other words, if we utilize the existing DSPs for 240 multipliers, we will save 23.09% of the LUTs for other modules.**

DATA_WIDTH	mul_dsp		mul_lut	
	LUT	DSP	LUT	DSP
4	23	0	23	0
6	36	0	36	0
8	61	0	61	0
9	88	0	88	0
10	102	0	102	0
11	0	1	132	0
12	0	1	148	0
16	0	1	265	0
24	0	2	595	0
32	47	4	1027	0

Let's make a **multiplier**.

Step 1: Generate a DSP from IP catalog

Step 2: Make "mul.v"

- Inputs: clk, w[7:0], x[7:0]
- Output: y[15:0]

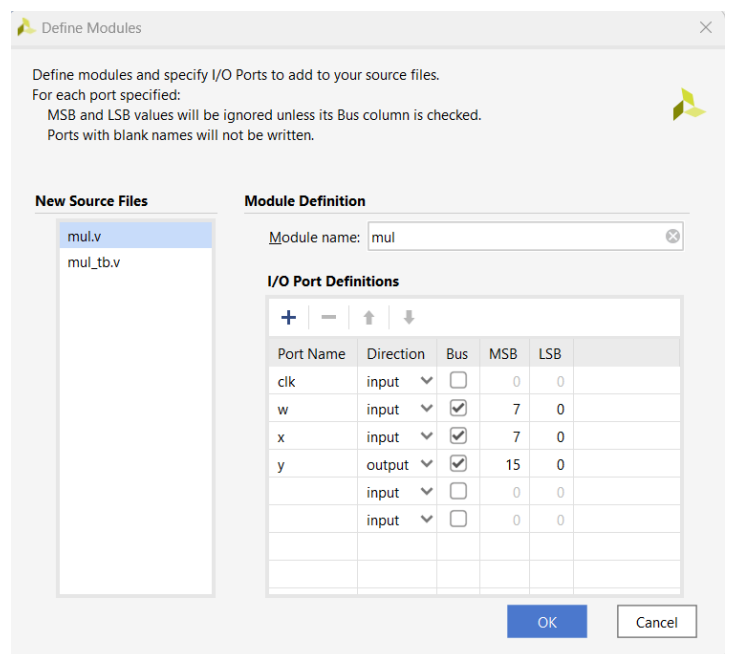
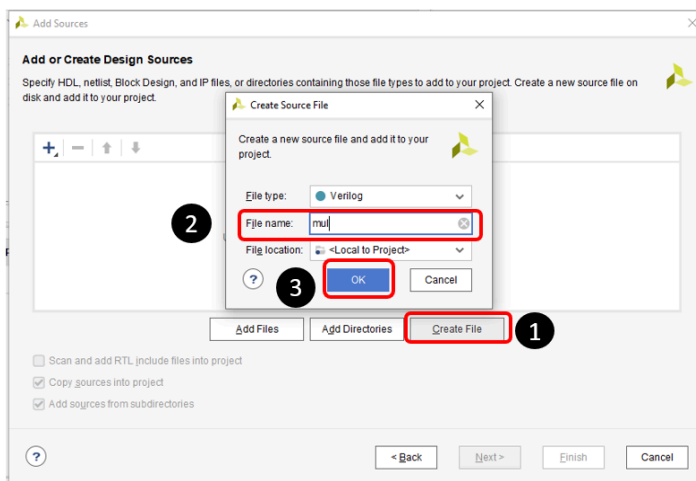
Step 3: Make inputs and outputs for the DSP

Step 4: Call a DSP instance

(This time, you don't have to use DSP generator because you can just use the one you made before. So, all you need to do here is to create a new code to instantiate the DSPs.)

To start with, "Add source">> Click "Create File" >> Make a file name "mul" in "Verilog" >> click "OK"

- Note: You can just download add the source code(mul.v)



You can make an **8-bit-by-8-bit multiplier** as follows.

```
mul.v
C:/Users/User/aix2022/aix2022.srscs/sources_1/new/mul.v

1 | `timescale 1ns / 1ps
2 | module mul(
3 |     input clk,
4 |     input [7:0] w,
5 |     input [7:0] x,
6 |     output [15:0] y
7 | );
8 |
9 |     wire [17:0] dsp_A, dsp_B;
10 |    wire [47:0] dsp_P;
11 |
12 |    assign dsp_A = w[7]? {10'b11_1111_1111, w} : {10'b00_0000_0000, w};
13 |    assign dsp_B = x[7]? {10'b11_1111_1111, x} : {10'b00_0000_0000, x};
14 |    assign y = dsp_P[15:0];
15 |
16 |    dsp_macro_0 u_dsp(.CLK(clk), .A(dsp_A), .B(dsp_B), .C(48'b0), .P(dsp_P));
17 | endmodule
18 |
```

Make inputs and outputs for the DSP

Negative Positive

2's complement for a negative number

An instance of DSP

Once you've created the design(mul.v), you need to make corresponding testbench(mul_tb.v) to simulate it. In general, clock generator, DUT(Design Under Test), and input stimulus generator is included in testbench.

Time scale: 1ns

```

1  `timescale 1ns / 1ps
2  module mul_tb;
3      reg clk;
4      reg rstn;
5      reg [7:0] w, x;
6      wire [15:0] y;
7      //
8      // DUT: multiplier
9      //
10     mul u_mul(
11         ./*input      */clk(clk),
12         ./*input [ 7:0] */w(w),
13         ./*input [ 7:0] */x(x),
14         ./*output[15:0] */y(y)
15     );
16
17     // Clock
18     parameter CLK_PERIOD = 10; //100MHz
19     initial begin
20         clk = 1'b1;
21         forever #(CLK_PERIOD/2) clk = ~clk;
22     end
                
```

Design under test:
multiplier (mul)

Clock: 100MHz
(Period = 10ns)

Registers (e.g. clk, w, x) are connected to the inputs
A wire (e.g., y) is connected to the output

A clock signal is defined in an "initial" block

- Initialized at 1
- State is changed at every 5ns

Initialized states
Set all registers to default values

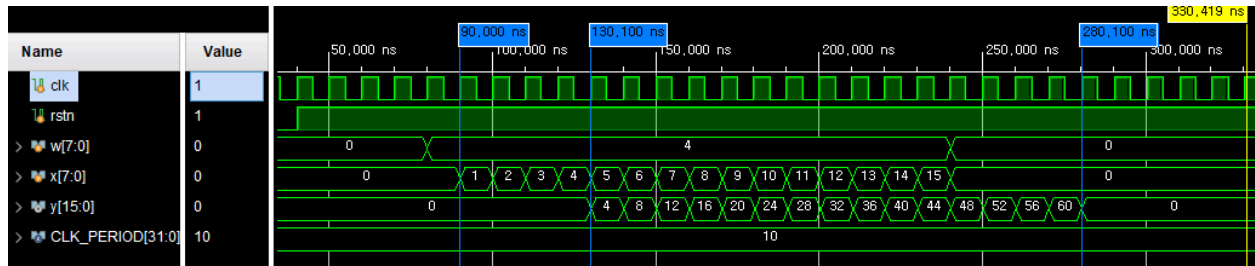
Generate test cases in 16 cycles
w: 4
x: 0 -> 1 -> 2 -> ... -> 15

```

23     integer i;
24     // Test cases
25     initial begin
26         rstn = 1'b0; // Reset, low active
27         w = 0;
28         x = 0;
29         i = 0;
30         #(4*CLK_PERIOD) rstn = 1'b1;
31
32         #(4*CLK_PERIOD)
33         for(i = 0; i<16; i=i+1) begin
34             @(posedge clk)
35                 w = 8'd4;
36                 x = i;
37         end
38
39         #(CLK_PERIOD)
40         @(posedge clk)
41             w = 8'd1;
42             x = 8'd1;
43     end
                
```

Now, you can run the simulation.

In the Tab "Simulation", click on "run simulation" >> "Run behavioral simulation".



Next, Let's make **an array of multiplier with MAC**(Multiplication and ACcumulation). You can download the source codes. They are for example MAC module and its testbench that calculates **(1)16 multiplications(mac.v)** and **(2)accumulations(adder_tree.v)** **every cycle**.

Assume that we compute the inner product of two vectors W and x , each with 16 elements:

$$y = \sum_{i=0}^{15} w_i * x_i$$

To do so, we'll instantiate sixteen multipliers that can execute sixteen multiplication, with an adder tree to add all the products. Inputs are 8 bit 2's complement numbers and the output is a 20 bit 2's complement number.

[mac.v]

- Input
 - clk, rstn Clock and reset. Active low
 - vld_i input valid signal
 - win 128 bits for weights → 16 eight-bit weights
 - din 128 bits for input pixels → 16 eight-bit pixels
- Output
 - vld_o Output valid signal
 - acc_o Accumulated output that has 20 bits.

```

1 timescale 1ns / 1ps
2
3 module mac(
4   input clk,
5   input rstn,
6   input vld_i,
7   input [127:0] win,
8   input [127:0] din,
9   output [19:0] acc_o,
10  output      vld_o
11 );
12

```

• Pseudo code

```

 $y_0^{(0)} = w_0 * x_0, \dots, y_{15}^{(0)} = w_{15} * x_{15}$  // N multipliers
 $y_0^{(1)} = y_0^{(0)} + y_1^{(0)}, \dots, y_7^{(1)} = y_{14}^{(0)} + y_{15}^{(0)}$  // N/2 adders
 $y_0^{(2)} = y_0^{(1)} + y_1^{(1)}, \dots, y_3^{(2)} = y_6^{(1)} + y_7^{(1)}$  // N/4 adders
 $y_0^{(3)} = y_0^{(2)} + y_1^{(2)}, \dots, y_1^{(3)} = y_2^{(2)} + y_3^{(2)}$  // N/4 adders
 $y_0^{(4)} = y_0^{(3)} + y_1^{(3)}$  // N/4 adders
 $Y = y_0^{(4)}$  // Output

```

[adder tree.v]

The outputs from the sixteen multipliers above are forwarded to an adder tree for accumulation. Use 'adder_tree.v' to compute a sum of 16 products.

To obtain an output, the adder tree uses **fifteen adders to perform hierarchical addition**.

At the first level, sixteen inputs are paired and then forwarded to **eight** adders that generate eight intermediate results. At the second level, the eight results are paired and moved to **four** adders that produce four results. Next, the four outputs are paired and inserted into **two** adders that output two results.

Finally, the results are forwarded to the **last** adder to generate the accumulated output.

□ $8+4+2+1=15$

```

13 // Internal signals
14
15
16 wire[15:0] y00;
17 wire[15:0] y01;
18 wire[15:0] y02;
19 wire[15:0] y03;
20 wire[15:0] y04;
21 wire[15:0] y05;
22 wire[15:0] y06;
23 wire[15:0] y07;
24 wire[15:0] y08;
25 wire[15:0] y09;
26 wire[15:0] y10;
27 wire[15:0] y11;
28 wire[15:0] y12;
29 wire[15:0] y13;
30 wire[15:0] y14;
31 wire[15:0] y15;

```



```

34 // 16 multipliers running in parallel
35
36
37 mul_u_mul_00(.clk(clk), .w(win[ 7: 0]), .x(din[ 7: 0]), .y(y00));
38 mul_u_mul_01(.clk(clk), .w(win[15: 8]), .x(din[15: 8]), .y(y01));
39 mul_u_mul_02(.clk(clk), .w(win[23:16]), .x(din[23:16]), .y(y02));
40 mul_u_mul_03(.clk(clk), .w(win[31:24]), .x(din[31:24]), .y(y03));
41 mul_u_mul_04(.clk(clk), .w(win[39:32]), .x(din[39:32]), .y(y04));
42 mul_u_mul_05(.clk(clk), .w(win[47:40]), .x(din[47:40]), .y(y05));
43 mul_u_mul_06(.clk(clk), .w(win[55:48]), .x(din[55:48]), .y(y06));
44 mul_u_mul_07(.clk(clk), .w(win[63:56]), .x(din[63:56]), .y(y07));
45 mul_u_mul_08(.clk(clk), .w(win[71:64]), .x(din[71:64]), .y(y08));
46 mul_u_mul_09(.clk(clk), .w(win[79:72]), .x(din[79:72]), .y(y09));
47 mul_u_mul_10(.clk(clk), .w(win[87:80]), .x(din[87:80]), .y(y10));
48 mul_u_mul_11(.clk(clk), .w(win[95:88]), .x(din[95:88]), .y(y11));
49 mul_u_mul_12(.clk(clk), .w(win[103:96]), .x(din[103:96]), .y(y12));
50 mul_u_mul_13(.clk(clk), .w(win[111:104]), .x(din[111:104]), .y(y13));
51 mul_u_mul_14(.clk(clk), .w(win[119:112]), .x(din[119:112]), .y(y14));
52 mul_u_mul_15(.clk(clk), .w(win[127:120]), .x(din[127:120]), .y(y15));

```

$y00 = w_0 * x_0$
 $win[7:0] = w_0$
 $din[7:0] = x_0$

$y15 = w_{15} * x_{15}$
 $win[127:120] = w_{15}$
 $din[127:120] = x_{15}$

To generate the output valid signal, we use a **delay line**. For example, it takes **four cycles to compute a multiplication** using the default DSP, while it takes **four cycles (=log2(16)) to accumulate** the sixteen multiplication results to generate an output.

□ Therefore, it takes **eight cycles of latency** to calculate $y = \sum_{i=0}^{15} w_i * x_i$.

- vld_i: when high, this indicates that 'win' and 'din' are valid inputs, and must be calculated
 - vld_d(delay): an internal variable to count the delay for vld_o to be high
 - vld_o(delay): when high, indicates that output acc_o is a valid calculation result
- Note that the eight cycles are the latency **between input and output**. If we insert multiple inputs consecutively, it will generate the consecutive outputs too, accordingly.

Now that you created the module design, you can simulate the module by creating testbench code(mac_tb.v). You can download and add it to your project.

Time scale: 1ns

Design under test:
mac

```

1 | `timescale 1ns / 1ps
2 |
3 | module mac_tb;
4 |   reg clk;
5 |   reg rstn;
6 |   reg vld_i;
7 |   reg [127:0] win, din;
8 |   wire [19:0] acc_o;
9 |   wire vld_o;
10 |
11 | // -----
12 | // DUT: multiplier
13 | // -----
14 |   mac_u_mac(
15 |     /*input      */clk(clk),
16 |     /*input      */rstn(rstn),
17 |     /*input      */vld_i(vld_i),
18 |     /*input [127:0] */win(win),
19 |     /*input [127:0] */din(din),
20 |     /*output [19:0] */acc_o(acc_o),
21 |     /*output      */vld_o(vld_o)
22 |   );

```

Registers (e.g. clk, win, din) are connected to the inputs

Wires (e.g., acc_o, vld_o) are connected to the output

Clock: 100MHz

A clock signal is defined in an "initial" block

- Initialized at 1
- State is changed at every 5ns

```

24 | // Clock
25 | parameter CLK_PERIOD = 10; //100MHz
26 | initial begin
27 |   clk = 1'b1;
28 |   forever #(CLK_PERIOD/2) clk = ~clk;
29 | end

```

Initialized states
Set all registers to default values

Generate test cases in 16 cycles
w: 4
x: 0 -> 1 -> 2 -> ... -> 15

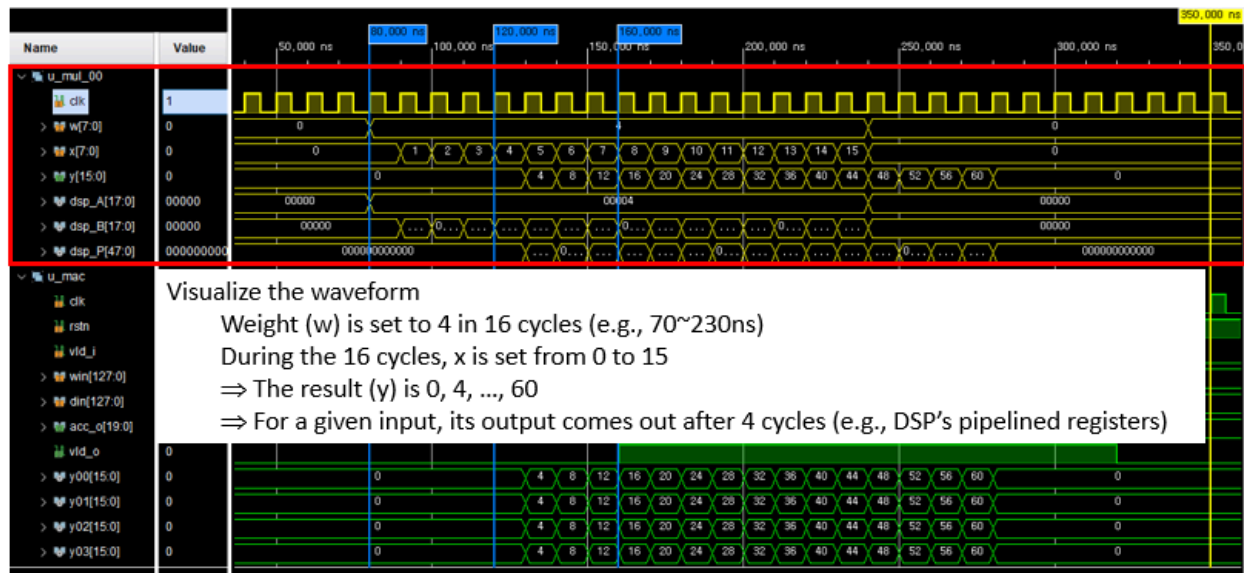
```

31 // Test cases
32 initial begin
33     rstn = 1'b0; // Reset, low active
34     vid_i = 0;
35     win = 0;
36     din = 0;
37     i = 0;
38     #(4*CLK_PERIOD) rstn = 1'b1;
39
40     #(4*CLK_PERIOD)
41     for(i = 0; i < 16; i=i+1) begin
42         @(posedge clk)
43             vid_i = 1'b1;
44             win = {16{8'd4}};
45             din[ 7: 0] = i;
46             din[15: 8] = i;
47             din[23:16] = i;
48             din[31:24] = i;
49             din[39:32] = i;
50             din[47:40] = i;
51             din[55:48] = i;
52             din[63:56] = i;
53             din[71:64] = i;
54             din[79:72] = i;
55             din[87:80] = i;
56             din[95:88] = i;
57             din[103:96] = i;
58             din[111:104] = i;
59             din[119:112] = i;
60             din[127:120] = i;
61     end

```

Note: we have 16 weights
and 16 inputs now.

When you run the simulation, you can see the waveform like this. You can compare multiplier and mac unit as follows:



3. Convolutional Kernel(Parallel MAC modules)

We observed that 16 multipliers and 15 adders could be used to calculate an inner product between two vectors with sixteen eight-bit elements. Recall that the AIX2024 SDK does an inference with a quantized model as follows:

- Do activation quantization and vectorize eight-bit input pixels (im2col_cpu_int8) and save them into **b**.
- Retrieve the quantized int8 weights(weights_int8) and save them into **a**.
- Do convolution, for example, inner product between **a(weight)** and **b(input activation)**, in parallel.

** Note that the loop “for (t =0; t <m; ++t)” iterates among m sets of the filters. In other words, each iteration produces one accumulated map.*

- Meanwhile, ‘gemm_nn_int8_int16’ takes two vectors from **a** and **b**, and calculates inner products as [mac.v] and the computed results are stored in **c**.

```
111 void forward_convolutional_layer_q(layer l, network_state state)
112 {
113
114     int out_h = (l.h + 2 * l.pad - l.size) / l.stride + 1; // output_height=input_height for stride=1 and pad=1
115     int out_w = (l.w + 2 * l.pad - l.size) / l.stride + 1; // output_width=input_width for stride=1 and pad=1
116     int i, j;
117     int const out_size = out_h*out_w;
118
119     typedef int16_t conv_t; // 1.output
120     conv_t *output_q = calloc(l.outputs, sizeof(conv_t));
121
122     state.input_int8 = (int8_t *)calloc(l.inputs, sizeof(int));
123     int z;
124     for (z = 0; z < l.inputs; ++z) {
125         int16_t src = state.input[z] * l.input_quant_multiplier;
126         state.input_int8[z] = max_abs(src, MAX_VAL_8);
127     }
128
129     // Convolution
130     int m = l.n;
131     int k = l.size*l.size*l.c;
132     int n = out_h*out_w;
133     int8_t *a = l.weights_int8;
134     int8_t *b = (int8_t *)state.workspace;
135     conv_t *c = output_q; // int16_t
136
137     // Use GEMM (as part of BLAS)
138     im2col_cpu_int8(state.input_int8, l.c, l.h, l.w, l.size, l.stride, l.pad, b);
139     int t; // multi-thread gemm
140     #pragma omp parallel for
141     for (t = 0; t < m; ++t) {
142         gemm_nn_int8_int16(1, n, k, 1, a + t*k, k, b, n, c + t*n, n);
143     }
144     free(state.input_int8);
145 }
```

Annotations:

- Eight input pixels Vectorize (points to lines 122-127)
- Eight quantized weights (weights_int8) (points to line 133)
- Do convolution in parallel (points to lines 140-143)

While **mac.v** does 16 multipliers in parallel, we can **stack multiple MAC modules** to replicate the loop “for (t = 0; t < m; ++t)” in the AIX2024 SDK.

For example, we **stack four MAC kernels (cnv_tb.v)**, with each module being responsible for one set of convolutional filters. In particular, there are four streams of weights **win[0:3]** that **apply for the same input pixel vector din** to produce four different accumulated outputs.

Read the hex file
into a buffer

Initialize four set of
convolutional filters

⇒ Come from a
quantized model

```

68 // Read the input file to memory
69 initial begin
70   $readmemh(INFILE, in_img, 0, FRAME_SIZE-1);
71 end
72
73 initial begin
74   rstn = 1'b0; // Reset, low active
75   vid_i = 0;
76   din = 0;
77   i = 0;
78
79 // CNN filters of four output channels
80 win[0][ 7: 0] = 8'd142; win[1][ 7: 0] = 8'd59; win[2][ 7: 0] = 8'd13; win[3][ 7: 0] = 8'd59;
81 win[0][15: 8] = 8'd151; win[1][15: 8] = 8'd181; win[2][15: 8] = 8'd244; win[3][15: 8] = 8'd135;
82 win[0][23:16] = 8'd215; win[1][23:16] = 8'd209; win[2][23:16] = 8'd255; win[3][23:16] = 8'd235;
83 win[0][31:24] = 8'd127; win[1][31:24] = 8'd19; win[2][31:24] = 8'd241; win[3][31:24] = 8'd128;
84 win[0][39:32] = 8'd163; win[1][39:32] = 8'd128; win[2][39:32] = 8'd127; win[3][39:32] = 8'd32;
85 win[0][47:40] = 8'd305; win[1][47:40] = 8'd95; win[2][47:40] = 8'd40; win[3][47:40] = 8'd90;
86 win[0][55:48] = 8'd229; win[1][55:48] = 8'd221; win[2][55:48] = 8'd252; win[3][55:48] = 8'd48;
87 win[0][63:56] = 8'd255; win[1][63:56] = 8'd121; win[2][63:56] = 8'd237; win[3][63:56] = 8'd52;
88 win[0][71:64] = 8'd113; win[1][71:64] = 8'd8; win[2][71:64] = 8'd1; win[3][71:64] = 8'd11;
89 win[0][79:72] = 8'd0; win[1][79:72] = 8'd0; win[2][79:72] = 8'd0; win[3][79:72] = 8'd0;
90 win[0][87:80] = 8'd0; win[1][87:80] = 8'd0; win[2][87:80] = 8'd0; win[3][87:80] = 8'd0;
91 win[0][95:88] = 8'd0; win[1][95:88] = 8'd0; win[2][95:88] = 8'd0; win[3][95:88] = 8'd0;
92 win[0][103:96] = 8'd0; win[1][103:96] = 8'd0; win[2][103:96] = 8'd0; win[3][103:96] = 8'd0;
93 win[0][111:104] = 8'd0; win[1][111:104] = 8'd0; win[2][111:104] = 8'd0; win[3][111:104] = 8'd0;
94 win[0][119:112] = 8'd0; win[1][119:112] = 8'd0; win[2][119:112] = 8'd0; win[3][119:112] = 8'd0;
95 win[0][127:120] = 8'd0; win[1][127:120] = 8'd0; win[2][127:120] = 8'd0; win[3][127:120] = 8'd0;

```

1	2a
2	45
3	5b
4	63
5	6a
6	6c
7	6f
8	6d

This time, use the testbench “**conv_tb**” to see how the convolution is parallelly processed with multiple macs and adder trees that we created before.

```

1 // timescale 1ns / 1ps
2
3 module cnv_tb;
4   parameter WIDTH = 128;
5   parameter HEIGHT = 128;
6   parameter INFILE = "../hex/butte";
7   localparam FRAME_SIZE = WIDTH * HEIGHT;
8   localparam FRAME_SIZE_W = $clog2(FRAME_SIZE);
9   reg [7:0] in_img [0:FRAME_SIZE-1];
10  reg clk;
11  reg rstn;
12  reg vid_i;
13  reg [127:0] win[0:3];
14  reg [127:0] din;
15  wire [19:0] acc_o[0:3];
16  wire vid_o[0:3];
17
18  mac u_mac_00(
19    .input(input),
20    .rstn(rstn),
21    .vid_i(vid_i),
22    .win(win[0]),
23    .din(din),
24    .output(acc_o[0]),
25    .vid_o(vid_o[0])
26  );
27
28  mac u_mac_01(
29    .input(input),
30    .rstn(rstn),
31    .vid_i(vid_i),
32    .win(win[1]),
33    .din(din),
34    .output(acc_o[1]),
35    .vid_o(vid_o[1])
36  );
37
38  mac u_mac_02(
39    .input(input),
40    .rstn(rstn),
41    .vid_i(vid_i),
42    .win(win[2]),
43    .din(din),
44    .output(acc_o[2]),
45    .vid_o(vid_o[2])
46  );
47
48  mac u_mac_03(
49    .input(input),
50    .rstn(rstn),
51    .vid_i(vid_i),
52    .win(win[3]),
53    .din(din),
54    .output(acc_o[3]),
55    .vid_o(vid_o[3])
56  );

```


- In the AIX2023 SDK, after convolution, it adds a bias, performs activation, and then descaling. The hardware module must replicate similar operations.

```

184 // Bias addition
185 int fil;
186 for (fil = 0; fil < l.n; ++fil) {
187     for (j = 0; j < out_size; ++j) {
188         output_q[fil*out_size + j] = output_q[fil*out_size + j] + l.biases_quant[fil];
189     }
190 }
191
192 // Activation
193 if (l.activation == RELU) {
194     for (i = 0; i < l.n * out_size; ++i) {
195         output_q[i] = (output_q[i] > 0) ? output_q[i] : 0;
196     }
197 }
198
199 // De-scaling or De-quantization
200 float ALPHA1 = 1 / (l.input_quant_multiplier * l.weights_quant_multiplier);
201 for (i = 0; i < l.outputs; ++i) { ... }
204
205 // Write data for the HW verification
206 //{{{
207 if (run_single_image_test) {
208     // Output Feature Map (OFM)
209     int z;
210     int next_input_quant_multiplier = 1;
211     for (z = state.index + 1; z < net.n; ++z) {
212         if (net.layers[z].type == CONVOLUTIONAL) {
213             next_input_quant_multiplier = net.layers[z].input_quant_multiplier;
214             break;
215         }
216     }
217     char file_output_femap[100];
218     sprintf(file_output_femap, sizeof(file_output_femap), "C:/skeleton/bin/log_femap/CONV%02d_output.hex", state.index);
219     FILE* fp = fopen(file_output_femap, "w");
220
221     // Data Format: [Channel, Width, Height]
222     for (int idx = 0; idx < out_size; idx++) { // OFM: Pixel index in ONE feature map
223         for (int chn = 0; chn < l.n; chn++) { // OFM: Channel/index of an feature map
224             int i = chn * out_size + idx; // OFM: Pixel index
225             uint8_t pixel = max_abs(l.output[i] * next_input_quant_multiplier, MAX_VAL_UINT_8);
226             fprintf(fp, "%02x\n", pixel);
227         }
228     }
229     if (fp) fclose(fp);
230 }
231 }
232 //}}}

```

Here is one example of the final design.



References

[1]. DSP48E1

<https://docs.xilinx.com/r/2021.1-English/ug1483-model-composer-sys-gen-user-guide/DSP48E1>

[2]. Deep Learning with INT8 Optimization on Xilinx Devices

<https://docs.xilinx.com/v/u/en-US/wp486-deep-learning-int8>